

What is an OS?

- OS is a software that sits between hardware and applications and provides interfaces to applications for accessing hardware

What is an application?

example.c

```
#include <stdio.h>
```

```
main() {
```

```
    printf("Hello world\n");
```

```
}
```

gcc example.c

./a.out

a.out is an application.

What is an application?

- Firefox
- Terminal
- Power point
- etc.

Hardware

- Disk
- Network device
- RAM
- CPU
- Monitor
- Keyboard
- etc.

Disk interface

512

0	1	2	3	4	5	6	7	8	9
---	---	---	---	---	---	---	---	---	---

```
write_to_sector(0)
write_to_sector(7)
read_from_sector(9)
```

The disk interface is the sequence of sectors. In this example, the sector size is 512 bytes. The disk interface allows programmers to read/write from/to sectors.

OS interface

```
fd = open(filename, "w");  
write(fd, buf, 32);
```

Even though the entire file is not stored in consecutive sectors, the OS interfaces allow you to access the file sequentially

OS interface is the file system. A file is divided into the contiguous blocks, where the block size is equal to sector size. Blocks can be stored in random sectors. The OS hides the complexity of blocks to sector translation from the programmers.

OS interface

```
fd = open(filename, "w");  
write(fd, buf, 32);
```

Even though the entire file is not stored in consecutive sectors, the OS interfaces allow you to access the file sequentially

BLOCK	SECTOR
0	15
1	23
2	43
3	29
...	...

The OS keeps the mapping from blocks to sectors.

What does an OS do?

- OS is a library
 - `write_to_console`, `write`, `read`, etc.

write_to_console

- An API that writes a character to the display device
- Applications are not allowed to use their own implementation of `write_to_console`
 - Different from the standard library where you can use a custom implementation of library functions, e.g., `strcpy`, `strcmp`, etc.

What does an OS do?

- OS is a library
 - `write_to_console`
- Enforce use of OS library
 - Access to hardware resources is only permitted through the OS library APIs

We can not let the applications access the console directly, because in that case, an application can hijack the console and prevent other applications from using them.

Isolation

- An application can't see the data of other applications
- Applications are untrusted
 - A gaming application can't see your password in a banking application
- OS is trusted
 - If our OS is malicious, then it can allow malicious applications to steal our passwords
 - We all trust our OS to protect us from malicious software

What does an OS do?

- OS is a library
 - `write_to_console`
- Enforce use of OS library
 - Access to hardware resources is only permitted through the OS library APIs
- Isolation

Sharing of hardware resources

- E.g., OS decides which application is going to use the display device at a given time

What does an OS do?

- OS is a library
 - `write_to_console`
- Enforce use of OS library
 - Access to hardware resources is only permitted through the OS library APIs
- Isolation
- Share hardware resources among applications

Tentative weekly lecture plan

1. Introduction to OS and x86 architecture
2. Threads
3. Synchronization
4. Processes, System calls
- 5, 6. Segmentation
- 7, 8. Paging
- 9, 10. File system
- 11, 12. Concurrency
13. Other OS designs

Books

- xv6 book
- xv6 code listing
- Operating System Concepts, 9th Edition, Wiley by Silberschatz, Galvin, Gagne
- Intel software developers manual (for x86 architecture)
- Pintos documentation

Exams

- You have to bring the xv6 book and xv6 code listing during the exams

Prerequisite

- Good C programming skills

Evaluation

- Programming : 20
- Homework : 10
- Mid semester : 25
- End semester : 40
- Refresher module : 5

Passing criteria

- Meet the following conditions
 - At least 15 marks in midsem + endsem (65)
 - At least 10 marks in other components (35)

Office hours

- Mon, Wed : You can meet me just after class
- Tue, Thurs : 3pm -3:30pm (A505, New academic)

Plagiarism

<https://www.iiitd.ac.in/academics/resources/academic-dishonesty>

What is inside a.out?

- `gcc -m32 -fno-pic example.c`
- `objdump -dx a.out | less`

a.out contains CPU (x86) instructions
Instructions are stored in memory

Physical address space

- Each byte in the main memory (RAM) has an address
- Using 32 bits, we can generate 2^{32} unique values
- In x86 32-bit architecture, the physical address space contains 2^{32} addresses

Physical address space

32-bit memory mapped devices	0xFFFFFFFF (4 GB)
unused	
	depends on RAM
Extended memory	
BIOS ROM	0x10000 (1 MB)
16-bit devices	0x0F000 (960 KB)
VGA display	0x0C000 (768 KB)
Low memory	0x0A000 (768 KB)
	0x00000

For larger RAM (> 4GB), PAE support is available in x86. PAE can only be used with page-tables (will discuss later). PAE allows a 48-bit physical address on 32-bit x86 architecture.

Registers

- A general-purpose register can hold a 32-bit value (data/address)
- X86 has eight general-purpose registers
 - EAX, ECX, EDX, EBX, ESI, EDI, EBP, ESP

Registers

31	15	7	0
	AH	AL	
	BH	BL	
	CH	CL	
	DH	DL	
	SI		
	DI		
	BP		
	SP		

16 bit	32 bit
AX	EAX
BX	EBX
CX	ECX
DX	EDX
	ESI
	EDI
	EBP
	ESP

What is inside a.out?

- `gcc -m32 -fno-pic example.c`
- `objdump -dx a.out |less`

a.out contains CPU (x86) instructions

Instructions are stored in memory

CPU reads the instructions from memory and executes them

x86 CPU interface

- The instruction pointer (EIP register) contains the memory address where the next instruction is stored

```
while (1) {  
    instruction = fetch(EIP);  
    execute instruction;  
    if (instruction doesn't modify EIP) {  
        EIP = EIP + length(instruction);  
    }  
}
```

x86 instructions

- Branch instructions can modify EIP
 - call, ret, jmp, je, jne, etc. are branch instructions
 - Transfer the control to other parts of the code
- Instructions can read/write from/to physical address space
- Instructions can perform arithmetic or logical operations on registers and physical addresses

x86 instructions

add \$1, %eax

An instruction consists of an opcode followed by one or more operands

Operands are constants, registers, direct memory addresses, and indirect memory addresses

x86 instructions

`add $1, %eax`

The instructions are the sequence of bytes

The above notation is used by the programmers to write an assembly program that can be converted into the bytecode using an assembler

Register mode

```
movl %eax, %edx
```

`%edx = %eax`

b,w,l denotes the size of operands

AT&T (gcc) syntax: opcode src, dst

Register mode

```
movb %al, %bl
```

`%bl = %al`

b,w,l denotes the size of operands

AT&T (gcc) syntax: opcode src, dst

Register mode

```
movw %ax, %bx
```

```
%bx = %ax
```

b,w,l denotes the size of operands

AT&T (gcc) syntax: opcode src, dst

Immediate

```
movl $0x123, %eax  
%eax = 0x123;
```

Direct memory access

```
movl 0x123, %edx
```

```
%edx = *((int32*)0x123);
```

(memory read)

Direct memory access

```
movl %edx, 0x123
```

```
*((int32*)0x123) = %edx;
```

(write to memory)

Indirect memory operand

`disp(base, index, scale)`

disp: 32 bit signed integer

base, index : register

scale: 1,2,4,8

address: $\text{disp} + \text{base} + (\text{index} * \text{scale})$

Indirect memory access

```
movl (%ebx), %edx
```

```
%edx = *((int32*)%ebx)
```

(memory read)

Indirect memory access

```
movl 4(%ebx), %edx
```

```
%edx = *((int32*)(%ebx+4))
```


Indirect memory access

```
movl (%ebx,%ecx,4), %edx
```

```
%edx = ((int32*)(%ebx))[%ecx]
```

$ebx[ecx]$

$ebx + (ecx * 4)$



```
int a[10];
```

```
a[2] = 100;
```

$a + (2 * 4)$

Indirect memory access

```
movl (%ebx,%ecx,1), %edx
```

ebx + ecx

```
%edx = *((int32*)(%ebx + %ecx))
```

Indirect memory access

```
movl %edx, (%ebx,%ecx,1)
```

```
*(int32*)(%ebx + %ecx) = %edx
```

(write to memory)

Indirect memory access

```
movl $0x0, (%ebx,%ecx,1)
```

```
*(int32*)(%ebx + %ecx) = 0x0
```

do we need to specify the “l” suffix

Indirect memory access

```
mov $0x0, (%ebx,%ecx,1)
```

is ambiguous

It is not evident from this instruction that, whether we want to store 1,2, or 4-byte value (0) into the destination memory address. A “b”, “w”, or “l” suffix is required to eliminate this ambiguity.

Indirect memory access

```
mov %eax, (%ebx,%ecx,1)
```

is not ambiguous

Here, we know that %eax contains a 4-byte value.