

## Data movement instructions

READ	WRITE	ADDRESS
mov 0x123, %eax	mov %eax, 0x123	0x123
mov (%ebx), %eax	mov %eax, (%ebx)	%ebx
mov ( <u>%ebx, %edx, 4</u> ), %eax	mov %eax, (%ebx, %edx, 4)	<u>%ebx + %edx * 4</u>
mov ( <u>, %edx, 4</u> ), %eax	mov %eax, (, %edx, 4)	<u>%edx * 4</u>
mov 0x100(%ebx, %edx, 4), %eax	mov %eax, 0x100(%ebx, %edx, 4)	0x100 + %ebx + (%edx * 4)

## jmp instruction

`jmp rel32`                    `// unconditional jump`

$EIP = EIP + rel32$

rel32 is 32 bit signed integer

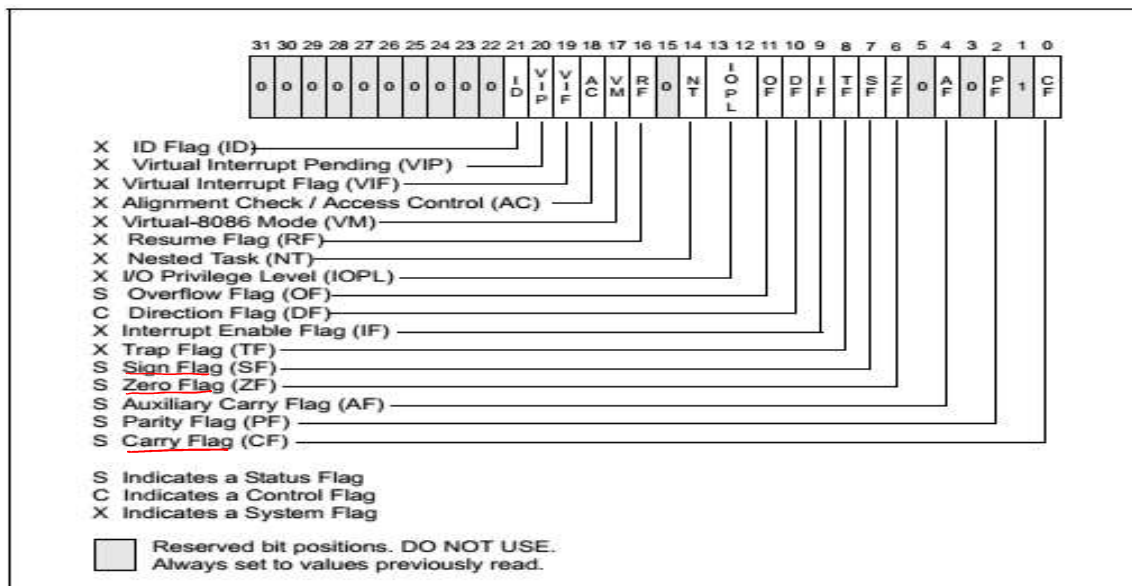
C equivalent:

`goto label;`

jmp instruction is an unconditional jump. This instruction sets the EIP to  $EIP + rel32$ .

## x86 EFLAGS register

- Section-3.4.3 in Intel's manual-1 (backpack)



**Figure 3-8. EFLAGS Register**

# EFLAGS

- Some instructions set certain bits in the EFLAGS register
  - Whether the last arithmetic instruction overflowed (OF)
  - Was positive/negative (SF)
  - Was zero / not zero (ZF)
  - Carry/borrow on add/subtract (CF)

# Compare

```
cmp %eax, %edx
```

Perform %edx - %eax

Sets the flags in the EFLAGS register accordingly

If %edx and %eax are equal the ZF will be set

If %edx > %eax then CF will not be set

if %edx < %eax the CF (borrow) will be set

jcc – jump if condition is met

je rel32                      // jump on equal

if (last operation was equal)    //(ZF = 1)

    EIP = EIP + rel32

If the ZF is set, then  $\text{src2} = \text{src1}$ . The x86 processor looks at the eflags register to find the value of zero flag(ZF). If ZF is set, then the x86 processor sets the EIP to the target address ( $\text{EIP} + \text{rel32}$ ).



jcc – jump if condition is met

ja rel32                      // jump on above

if (in last operation  $\text{src2} > \text{src1}$ )              // CF = 0 and ZF = 0  
    EIP = EIP + rel32

If the ZF is set then  $\text{src2} = \text{src1}$ . ja jumps to the target address (EIP + rel32) only if  $\text{src2} > \text{src1}$ . The x86 processor looks at the eflags register to find the value of the carry flag(CF) and the zero flag(ZF). If both of them are not set, then it sets the EIP to the target address (EIP + rel32).

jcc – jump if condition is met

jb rel32                      // jump on below

if (in last operation  $\text{src2} < \text{src1}$ )              // CF = 1

    EIP = EIP + rel32

jcc – jump if condition is met

jae rel32                      // jump on above or equal

if (in last operation  $\text{src2} \geq \text{src1}$ )    // CF = 0 or ZF = 1

    EIP = EIP + rel32

## Example

```
if (%eax < %edx) {  
}
```

## Example

```
if (%eax < %edx) {  
}
```

```
cmp %eax, %edx
```

```
jbe label
```

```
// jump if %edx <= %eax
```

```
// if body
```

```
label:
```

## Example

```
if (%eax < %edx) { //if body }  
else { //else body }
```

```
cmp %eax, %edx  
jbe label1 // jump if %edx <= %eax  
// if body  
jmp label2:  
label1:  
// else body  
label2:
```

After if body, we need to insert an unconditional jump to skip the else body.

## Addition

```
add %eax, %edx
```

```
%edx = %edx + %eax
```

## Addition

add %eax, (%edx)

-1·edx

$*((\text{int}32*)\%edx) = *((\text{int}32*)\%edx) + \%eax;$



## Logical AND

`and %eax, %edx`

`%edx = %edx & %eax;`

## Logical AND

`and %eax, (%edx)`

`*((int32*)%edx) = %edx & *((int32*)%eax);`

## References (available on backpack)

- A Guide to Programming Intel IA32 PC Architecture
- Intel manual - 2

## Next homework

- Read Intel manual-2 to find the meaning of some x86 instructions
- The Intel syntax is different from AT&T syntax
  - In Intel's syntax, the first operand is the destination
- Submit your handwritten homework in the submission box placed at the old academic building (2<sup>nd</sup> floor)

## Local variables

```
int foo(int x, int y)
{
    int v1, v2;
    int arr[64];
    ....
    bar(v1, v2);
}
```

Let us come back to our discussion on applications. How do local variables are allocated and deallocated in a C program?

## Local variables

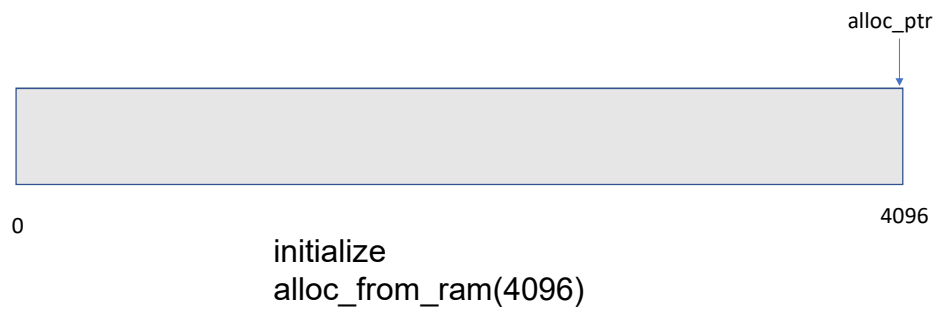
- Local variables are allocated and destroyed automatically by the compiler
- A straightforward strategy is to allocate all the local variables on function entry and deallocate them just before the function exit

## Local variable

```
void bar() {  
    int x1, x2, x3;  
    ...  
}
```

```
void foo() {  
    int v1, v2, v3;  
    bar();  
    ...  
}
```

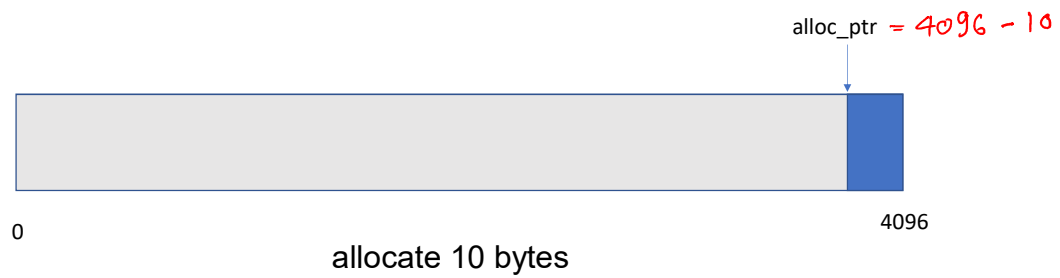
## Bump allocator



Initially, bump allocator allocates a contiguous memory area from the OS. It maintains an `alloc_ptr` that is set to the end of the memory area.




## Bump allocator




On every allocation, the `alloc_ptr` is advanced in the reverse direction by the size of the allocation. The new value of the `alloc_ptr` is returned to the user of the bump allocator.


## Bump allocator



## Bump allocator




## Bump allocator




During deallocation, the bump allocator advances the `alloc_ptr` by the size of the deallocation.

## Bump allocator



# Bump allocator



```
void initialize() {  
    char *ptr = alloc_from_ram(4096);  
    alloc_ptr = ptr + 4096;  
}
```

```
char *alloc(size_t size) {  
    alloc_ptr -= size;  
    return alloc_ptr;  
}
```

```
void dealloc(size_t size) {  
    alloc_ptr += size;  
}
```

## OS maintains a simple allocator

- `alloc_from_ram(size_t size)`
  - applications can use this interface to allocate space from RAM of size multiples of 4096
  - OS divides RAM into chunks of 4096 bytes called pages
  - OS maintains a list of free pages
  - Similar to `alloc_pages` in Linux kernel

## Physical address space

<b>32-bit memory mapped devices</b>	0xFFFFFFFF (4 GB)
<b>unused</b>	
	depends on RAM
<b>Extended memory</b>	
<b>BIOS ROM</b>	0x10000 (1 MB)
<b>16-bit devices</b>	0x0F000 (960 KB)
<b>VGA display</b>	0x0C000 (768 KB)
<b>Low memory</b>	0x0A000 (768 KB)
	0x00000

OS divides the entire physical space into the contiguous chunks of 4096 bytes. These chunks are also called pages. We can allocate contiguous pages (i.e., size multiples of 4096) using `alloc_from_ram` API.



## Local variable allocator

```
main () {  
    int a, b, c;   
    foo();  
}  
foo() {  
    int x, y;  
    bar();  
}  
bar () {  
    int a, b, c, x, y;  
}
```

### Bump allocator

```
void initialize();  
char *alloc(size_t size);  
void dealloc(size_t size);
```

On function entry, compiler insert calls to allocator to allocate the local variables. Just before returning from a function, the compiler deallocates space allocated for local variables. A bump allocator can be used for allocating and deallocating local variables. Bump allocator requires that the memory deallocation happens in the reverse order of their allocation that is true in this case.

## Local variable allocator

```
main () {           // initialize
    int a, b, c;     // alloc(12)
    foo();
}                   // dealloc(12)
foo() {
    int x, y;        // alloc(8)
    bar();
}                   // dealloc(8)
bar () {
    int a, b, c, x, y; // alloc(20)
}                   // dealloc(20)
```

### Bump allocator

```
void initialize();
char *alloc(size_t size);
void dealloc(size_t size);
```

## Local variable allocator

<pre>main () {     int a, b, c;     foo(); }</pre>	<pre>// initialize // alloc(12)  // dealloc(12)</pre>	<b>During execution</b>
<pre>foo() {     int x, y;     bar(); }</pre>	<pre>// alloc(8)  // dealloc(8)</pre>	<pre>initialize - alloc(12) - alloc(8) - alloc(20) - dealloc(20) - dealloc(8) - dealloc(12) -</pre>
<pre>bar () {     int a, b, c, x, y; }</pre>	<pre>// alloc(20) // dealloc(20)</pre>	

The alloc and dealloc routines are called in this order during the execution of this program.

## Parameter passing

```
foo() {  
    int x, y;  
    bar(x, y);  
}  
bar(int a1, int a2) {  
    return a1 + a2;  
}
```

### Bump allocator

```
void initialize();  
char *alloc(size_t size);  
void dealloc(size_t size);
```

Parameters are similar to local variables, except their values are initialized by the caller. For example, before calling bar foo routine sets the values of a1 and a2 to x and y, respectively. It is the caller's (foo in this case) responsibility to allocate space for the parameters and initialize them.

## Parameter passing

```
foo() {  
  int x, y;           // alloc(8)  
  bar(x);             → alloc(4)  
                      → dealloc(4)  
}  
                        // dealloc(8)  
bar(int x) {  
  int a, b;           // alloc(8)  
  baz(a);             → alloc(4)  
                      → dealloc(4)  
}  
                        // dealloc(8)  
baz(int y) {  
  int x, y, a, b;     // alloc(16)  
}  
                        // dealloc(16)
```

### During execution

```
alloc(8)  
alloc(4)  
dealloc(4)  
dealloc(8)  
alloc(8)  
alloc(4)  
alloc(16)  
dealloc(16)  
  
dealloc(4)  
dealloc(8)  
dealloc(4)  
dealloc(8)
```

Before every function call, the compiler can call bump allocator APIs to allocate space for parameters and deallocate them after the target function returns.

## Bump allocator

- To implement the bump allocator, we need to store the `alloc_ptr`
- Compilers reserve `%esp` register to store the `alloc_ptr`
  - `%esp` is called the stack pointer
  - RAM space allocated by the bump allocator is called the stack

## Calling convention

- Compiler compiles one routine at a time
- Often a routine doesn't know who are the potential callers
  - e.g., printf
- The caller or callee must follow some convention for the parameter passing
  - gcc 32-bit compiler passes parameters via stack

## Return address

```
foo() {  
    baz();  
}  
bar(){  
    baz();  
}  
baz(){  
}
```

How does baz know where to return (foo or bar). Along with the parameters, the return address is also passed by the caller.



# Calling convention

- The contract between caller and callee in x86:
  - On entry to a function call
    - %eip points to the first instruction of a function
    - %esp points to the return address
    - %esp+4 points to the first argument
    - %esp+8 points to the second argument
    - and so on

# Calling convention

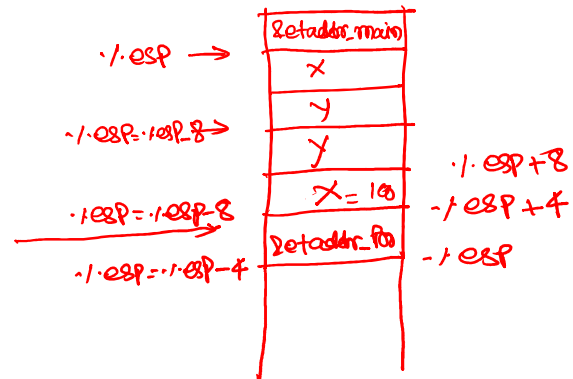
- The contract between caller and callee in x86:
  - After the callee returns
    - %eip points to the return address
    - %esp points to the argument pushed by the caller
    - called function may have trashed arguments
    - %eax (and %edx if the return type is 64 bit) contains the return value

## Example

```
void foo() {
    int x = 10;
    int y = 20;
    bar(x, y);
}
```

```
int bar(int x, int y) {
    return x + y;
}
```

$\rightarrow x = 100$



On entry to foo function stack pointer (%esp) points to the return address in the caller (say main). On entry to bar function, %esp points to return address in foo, (%esp + 4) points to x (first argument), (%esp + 8) points to y (second argument). After returning from bar %esp points to the argument pushed by foo, and the %eax register contains the return value of the bar routine.

## Push

```
push %eax
```

```
sub $4, %esp
```

```
mov %eax, (%esp)
```

- allocates 4-bytes on the stack and initializes them with the value of the %eax register

# Pop

```
pop %eax
```

```
mov (%esp), %eax
```

```
add $4, %esp
```

- deallocates 4-bytes from the stack and returns the value of the deallocated memory in the %eax register

## call instruction

call rel32

retaddr = EIP + length(EIP)

EIP = EIP + rel32 //rel32 is 32 bit signed integer

push retaddr

C equivalent:

bar();

call instruction set the EIP to EIP + rel32. In addition to this, call instruction also pushes the return address (retaddr) on the stack. x86 instructions are complex instructions and can do multiple tasks in one instruction.

## ret instruction

ret

pop %eip

C equivalent:


return

ret instruction pops the return address from the top of the stack and set the %eip to popped value. There is no such instruction pop %eip. The semantics of ret instruction is similar to the possible semantics of “pop %eip” (if it would have existed).

## Example

```
void foo() {
    int x = 10;
    int y = 20;
    bar(x, y);
}

int bar(int x, int y) {
    return x + y;
}
```



```
bar:
    mov 4(-1.ESP), %eax
    mov 8(-1.ESP), %ecx
    add -1.ecx, %eax
    ret
```

```
foo:
    sub $8, -1.ESP
    mov $10, 4(-1.ESP)
    mov $20, 8(-1.ESP)
    sub $8, -1.ESP
    mov 12(-1.ESP), %eax
    mov -1.eax, (-1.ESP)
    mov 8(-1.ESP), %eax
    mov -1.eax, 4(-1.ESP)
    call bar
    add $8, -1.ESP
    add $8, -1.ESP
    ret
```


This slide shows one possible code generation that follows the gcc calling convention.



## Example

```
void foo() {
    int x = 10;
    int y = 20;
    bar(&x, &y);
}

int bar(int *a, int *b) {
    return (*a) + (*b);
}
```



baz:

```
mov 4(%esp), %eax
mov (%eax), %ecx
mov 8(%esp), %eax
mov (%eax), %edx
add %ecx, %edx
mov %edx, %eax
ret
```

foo:

```
sub $8, %esp
mov $10, 4(%esp)
mov $20, 8(%esp)
sub $8, %esp
mov %esp, %eax
add $12, %eax
mov %eax, (%esp)
sub $4, %eax
mov %esp, %eax
add $8, %eax
mov %eax, 4(%esp)
call baz
add $8, %esp
add $8, %esp
ret
```

This is an example of pass by reference. Notice that in this case, we are passing the address of x and address of y to the bar routine. Because x and y are allocated on the stack, their addresses are stack addresses and are at a constant offset from %esp (known to the compiler).