


- Read chapter-0 from the xv6 book


before exec



before exec

32-bit memory mapped devices	0xFFFFFFFF (4 GB)
unused	
	depends on RAM
<del>shell</del>	
Extended memory	
BIOS ROM	0x10000 (1 MB)
16-bit devices	0x0F000 (960 KB)
VGA display	0x0C000 (768 KB)
	0x0A000 (768 KB)
Low memory	0x00000


after exec



after exec

32-bit memory mapped devices	0xFFFFFFFF (4 GB)
unused	
	depends on RAM
<del>IS</del>	
Extended memory	
BIOS ROM	0x10000 (1 MB)
16-bit devices	0x0F000 (960 KB)
VGA display	0x0C000 (768 KB)
Low memory	0x0A000 (768 KB)
	0x00000

before fork




before fork

32-bit memory mapped devices	0xFFFFFFFF (4 GB)
unused	
	depends on RAM
<del>shell</del>	
Extended memory	
BIOS ROM	0x10000 (1 MB)
16-bit devices	0x0F000 (960 KB)
VGA display	0x0C000 (768 KB)
	0x0A000 (768 KB)
Low memory	0x00000



after fork




after fork

32-bit memory mapped devices	0xFFFFFFFF (4 GB)
unused	
	depends on RAM
<del>shell</del>	
Extended memory	
<del>shell</del>	
BIOS ROM	0x10000 (1 MB)
16-bit devices	0x0F000 (960 KB)
VGA display	0x0C000 (768 KB)
	0x0A000 (768 KB)
Low memory	0x00000

# fork

- fork system call creates a new process that is identical to the child process
  - new memory (heap+stack+code+global) is allocated for the child process
  - the entire memory of the parent process is copied to the child process
  - the register values of the parent process are copied to the corresponding registers in the child process
  - after the fork returns changing the value of a variable in the parent doesn't change the variable's value in child and vice versa


before exit



before exit

32-bit memory mapped devices	0xFFFFFFFF (4 GB)
unused	
	depends on RAM
<i>Shell</i>	
Extended memory	
BIOS ROM	0x10000 (1 MB)
16-bit devices	0x0F000 (960 KB)
VGA display	0x0C000 (768 KB)
Low memory	0x0A000 (768 KB)
	0x00000

after exit



after exit

32-bit memory mapped devices	0xFFFFFFFF (4 GB)
unused	
	depends on RAM
Extended memory	
BIOS ROM	0x10000 (1 MB)
16-bit devices	0x0F000 (960 KB)
VGA display	0x0C000 (768 KB)
Low memory	0x0A000 (768 KB)
	0x00000

## System calls

- exec
- fork
- exit
- wait
  - the parent waits until one of the children exits



shell


\$

user space



shell

kernel space



kernel

shell is waiting for  
the user input

shell


\$ ls

user space



shell

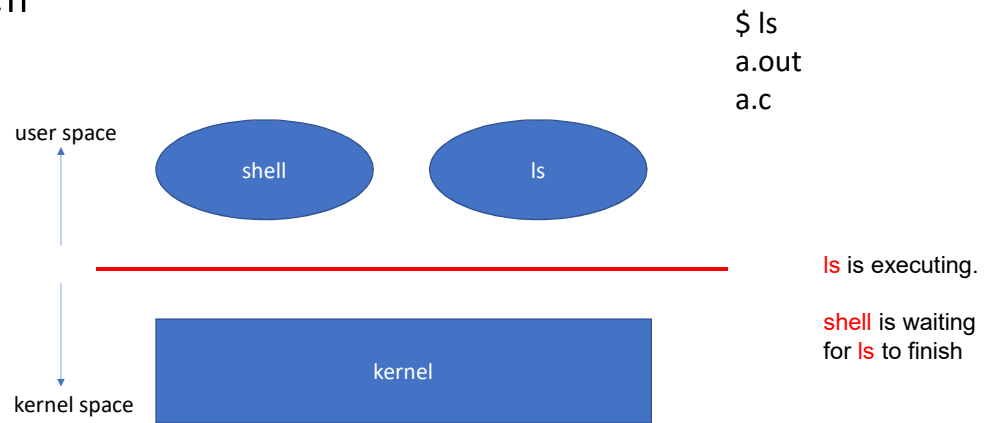
kernel space



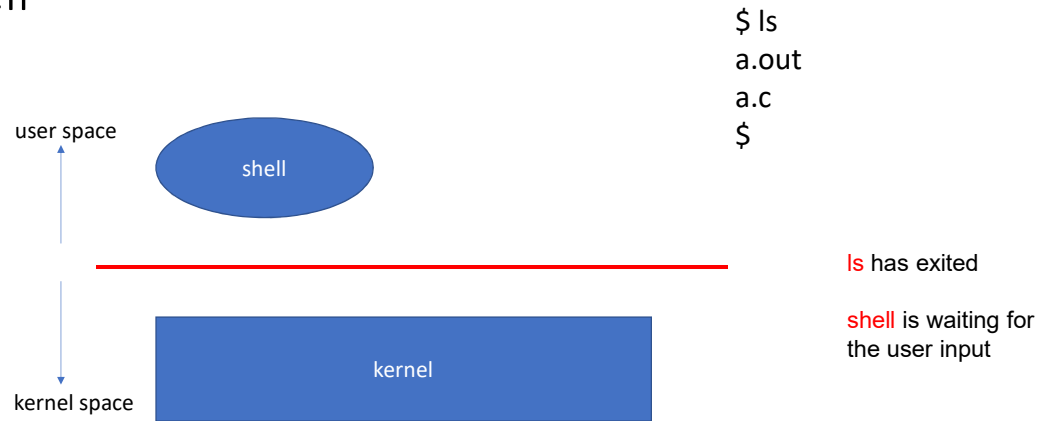
kernel

user entered **ls**

# shell



# shell



## shell

```
while (1) {  
    write(1, "$ ", 2);  
    read_command(0, cmd, args);  
    pid = fork();  
    if (pid == 0) {  
        exec(cmd, args);  
    } else if (pid > 0) {  
        wait();  
    } else  
        printf("Failed to fork\n");  
}
```

current working  
directory is copied to  
the child during fork.

**ls** program gives the  
same result in the  
parent and the child.

The child process does the exec system call to transform itself into an executable (in cmd) program. The parent (shell) process waits until the child (cmd) terminates.

## creat

- `int creat(const char *pathname, int flags);`
  - creates a new file
  - if the file already exists, truncates it
  - permissions (who can access the file) are set according to flags
  - On successful, an integer file descriptor is returned
  - the file descriptor can be used to write to file

## file descriptor

- file descriptor can refer to a file as well as a device
- file descriptor 0 refers to the input device (e.g., keyboard)
  - 0 points to standard input
- file descriptor 1, 2 refers to the output device (e.g., a display screen)
  - 1 points to standard output
  - 2 points to standard error

## write

- `ssize_t write(int fd, const void *buf, size_t count);`
  - write up to count bytes from buf to file referred by the file descriptor fd
  - returns the number of bytes that were written to the file
- To know more
  - type `man 2 write` in your terminal



## read

- `ssize_t read(int fd, void *buf, size_t count);`
  - read up to count bytes from the file referred by file descriptor fd into buf
  - returns the number of bytes read
  - If the fd is standard input, read system call waits for the user input
- To know more
  - type `man 2 read` in your terminal

## demo

- creat
- write
- write to console
- read from console


## File descriptor table

- A file descriptor is a small integer
- For every process, OS maintains a table of descriptors
  - open files, devices, etc.
- 0, 1, and 2 are reserved for standard input, standard output, and standard error
- The process may create/open new files using `creat` and `open` system calls
  - `creat` and `open` system calls add a new entry to the file descriptor table

## File descriptor table

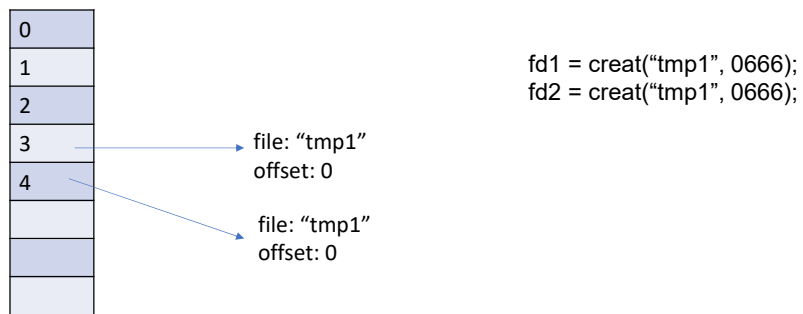
0	→ stdin
1	→ stdout
2	→ stderr

## File descriptor table



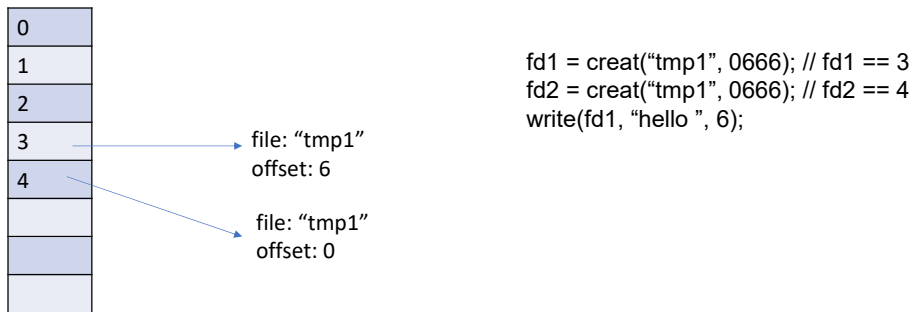
The create system call allocates an available descriptor in the file descriptor table. The file descriptor table entry points to a structure that contains an offset field. The offset is advanced by the number of bytes written during the write system call. The offset is also modified during the read if the file was opened for reading.

## File descriptor table



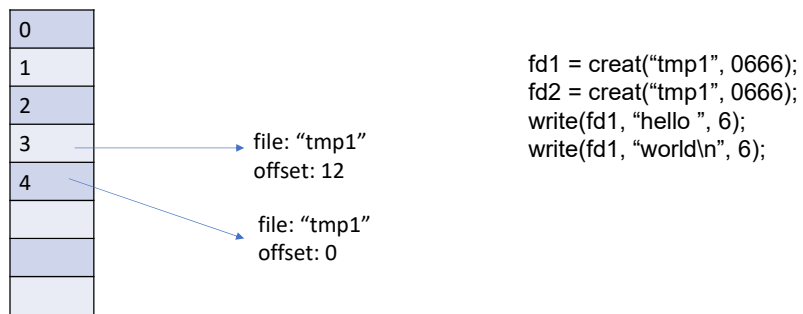
Two descriptors can point to the same file, but their offsets are different.

## File descriptor table



Write to fd1 updates the offset in the target pointer to struct at index 3 in the file descriptor table.

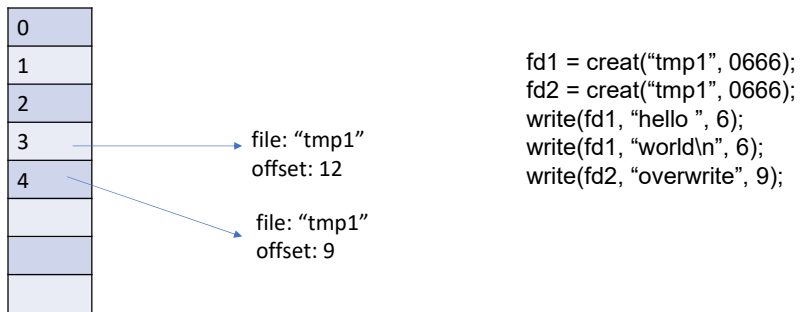
## File descriptor table



Write to fd1 updates the offset in the target pointer to struct at index 3 in the file descriptor table.



## File descriptor table



Write to fd2 updates the offset in the target pointer at index 4 in the file descriptor table. This write will overwrite the contents of the tmp1 file, starting at offset 0.

## I/O redirection

```
int main() {  
    write(1, "hello\n", 6);  
    return 0;  
}
```

```
a.out  
./a.out  
hello  
./a.out > tmp  
creates a new file tmp  
writes "hello\n" in tmp
```

## File descriptor table

0
1
2
3

file: "tmp1"  
offset: 0

```
fd1 = creat("tmp1", 0666);
```

creat always returns the lowest  
available descriptor in the file  
descriptor table

## File descriptor table

0
1
2

```
fd1 = creat("tmp1", 0666);
```

creat always returns the lowest available descriptor in the file descriptor table

```
close(fd1);
```

close system call removes the file descriptor fd1 from the file descriptor table

## File descriptor table

0
2

```
fd1 = creat("tmp1", 0666);  
creat always returns the lowest  
available descriptor in the file  
descriptor table
```

```
close(fd1);  
close system call removes the  
file descriptor fd1 from the file  
descriptor table
```

```
close(1);
```

## File descriptor table

0
1
2

→ file: "tmp1"  
offset: 0

```
fd1 = creat("tmp1", 0666);  
creat always returns the lowest  
available descriptor in the file  
descriptor table
```


```
close(fd1);  
close system call removes the  
file descriptor fd1 from the file  
descriptor table
```

```
close(1);  
fd1 = creat("tmp1", 0666);  
// fd1 is 1 at this point
```

## fork

- The child process inherits the file descriptor table from the parent
- The OS creates an identical file descriptor table for the child

## Before fork






## After fork




## After exec



Parent

file: "tmp1"  
offset: 0

On exec, the process  
retains file descriptors.



Child after exec

file: "tmp1"  
offset: 0

`./a.out > tmp1`

```
int main() {  
    write(1, "hello\n", 6);  
    return 0;  
}
```

```
a.out  
./a.out  
hello  
./a.out > tmp1  
creates a new file tmp1  
writes "hello\n" in tmp1
```

```
while (1) {  
    write(1, "$ ", 2);  
    read_command(0, cmd, args);  
    pid = fork();  
    if (pid == 0) {  
        close(1);  
        creat("tmp1", 0666);  
        exec(cmd, args);  
    } else {  
        wait();  
    } else  
        printf("Failed to fork\n");  
}
```

The two statements highlighted in red will make sure that the file descriptor 1 is now pointing to the tmp1 file. When the a.out program is loaded (after the exec system call), the file descriptor table remains the same as the child process. All writes to descriptor 1 will be written to tmp1.

`./a.out < tmp1`

```
int main() {  
    read(0, buf, 128);  
    write(1, buf, strlen(buf));  
    return 0;  
}
```

a.out  
./a.out  
waiting for input

`./a.out < tmp1`

```
int main() {  
    read(0, buf, 128);  
    write(1, buf, strlen(buf));  
    return 0;  
}
```

```
a.out  
./a.out  
waiting for input  
hello\n    // user enters  
writes "hello\n" to the console
```

`./a.out < tmp1`

```
int main() {  
    read(0, buf, 128);  
    write(1, buf, strlen(buf));  
    return 0;  
}
```

```
a.out  
./a.out  
waiting for input  
hello\n    // user enters  
writes "hello\n" to the console  
  
./a.out < tmp1  
writes contents of tmp1 to the console
```

`./a.out < tmp1`

```
int main() {  
    read(0, buf, 128);  
    write(1, buf, strlen(buf));  
    return 0;  
}
```

a.out  
./a.out  
waiting for input  
hello\n // user enters  
writes "hello\n" to the console

`./a.out < tmp1`  
writes "contents of tmp1" to the console

```
while (1) {  
    write(1, "$ ", 2);  
    read_command(0, cmd, args);  
    pid = fork();  
    if (pid == 0) {  
        close(0);  
        open("tmp1", O_RDONLY);  
        exec(cmd, args);  
    } else {  
        wait();  
    } else  
        printf("Failed to fork\n");  
}
```

This command redirects the contents of tmp1 to the stdin of a.out. Again using the trick highlighted in red, we can redirect the input from tmp1 to the standard input of a.out.

./a.out > tmp1 2>&1

```
int main() {  
    write(1, "hello ", 6);  
    write(2, "world\n", 6);  
    return 0;  
}
```


```
a.out  
./a.out  
hello world  
./a.out > tmp1 2>&1  
creates a new file tmp1  
writes "hello world\n" in tmp1
```

```
while (1) {  
    write(1, "$ ", 2);  
    read_command(0, cmd, args);  
    pid = fork();  
    if (pid == 0) {  
        close(1);  
        creat("tmp1", 0666);  
        close(2);  
        creat("tmp1", 0666);  
        exec(cmd, args);  
    } else {  
        wait();  
    } else  
        printf("Failed to fork\n");  
}
```

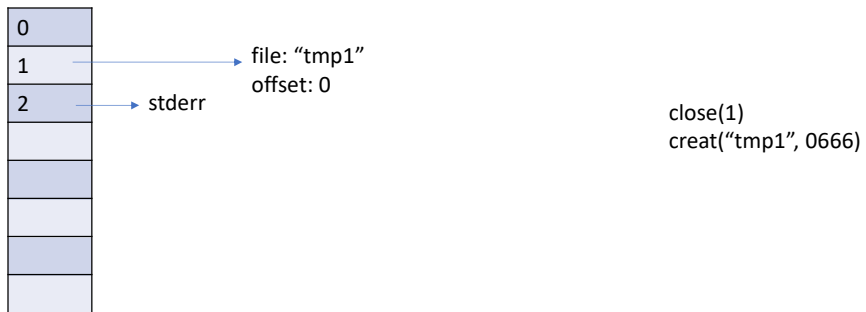
This command redirects the writes to file descriptors 1 and 2 to the tmp1 file. However, the code highlight in red is not suitable for this purpose. In this code, stdout (1) and stderr (2) will overwrite the contents of each other.



## File descriptor table



## File descriptor table



## File descriptor table

0
1
2

file: "tmp1"  
offset: 0 *6*

file: "tmp1"  
offset: 0 *6*

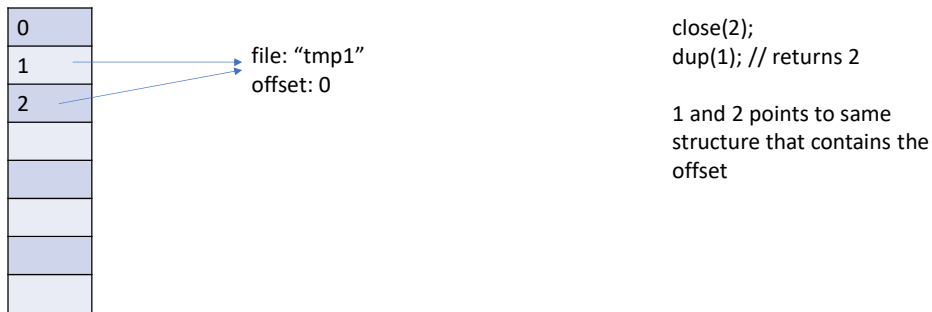
*"hello"*  
*"world"*

```
close(1)
creat("tmp1", 0666)
close(2)
creat("tmp1", 0666)
```

## dup

- `int dup(int oldfd);`
  - allocates a new descriptor and creates an alias of `oldfd` in the new file descriptor

## File descriptor table



The dup command takes a file descriptor as input and allocates a new file descriptor. The dup system call make sure that both input and new file descriptor are pointing to the same structure.

`./a.out > tmp1 2>&1`

```
int main() {  
    write(1, "hello ", 6);  
    write(2, "world\n", 6);  
    return 0;  
}  
  
a.out  
./a.out  
hello world  
./a.out > tmp1 2>&1  
create a new file tmp1  
write "hello world\n" in tmp1  
  
while (1) {  
    write(1, "$ ", 2);  
    read_command(0, cmd, args);  
    pid = fork();  
    if (pid == 0) {  
        close(1);  
        creat("tmp1", 0666);  
        close(2);  
        dup(1);  
        exec(cmd, args);  
    } else {  
        wait();  
    } else  
        printf("Failed to fork\n");  
}
```

The code highlighted in red is the correct solution.

## pipe

- pipe system call returns a pair of descriptor
- The first descriptor can be used for reading; the second descriptor can be used for writing

```
int p[2];  
pipe(p);  
p[0] → reading  
p[1] → writing
```

# pipe

```
[ int p[2];  
  pipe(p);
```





## pipe

- If the input file descriptor in the read or write system call is pipe
  - if no data is available, read system call waits until some data is written at the write end or all the file descriptors at the write end are closed
  - If the pipe is full, write system call waits until some data is consumed at the read end or all the file descriptors at the read end are closed
- read and write semantics for files are non-blocking

## pipe

```
int p[2];  
pipe(p);  
write(p[1], "hello", 5);  
read(p[0], buf, 5);  
// buf = "hello"
```

## Inter process communication (IPC)

```
int fd[2];  
pipe(fd);  
pid = fork();  
if (pid > 0)  
    write(fd[1], "hello", 5);    // parent sends "hello"  
else  
    read(fd[0], buf, 5);         // child receives "hello"
```

What happens if the child  
gets scheduled before the  
parent?

It will wait for the parent to  
do the write.

## pipe in shell

- demo

## pipe in shell

```
int main() {  
    return write(1, "hello\n", 6);  
}
```

a.out

```
int main() {  
    char buf[128];  
    read(0, buf, 128);  
    return write(1, buf, strlen(buf));  
}
```

b.out

## pipe in shell

`./a.out | ./b.out`



No need of temporary files  
file I/O is very slow  
`./a.out | ./b.out` is equivalent to  
`./a.out > tmp`  
`./b.out < tmp`

./a.out | ./b.out

no disk I/O!  
no deletion of temporary files  
no disk space needed  
parallel execution of pipeline stages  
no rewriting of applications  
faster IPC via memory.

```
int fd[2];
pipe(fd);
pid = fork();
if (pid == 0) {
    char *param[2] = {"a.out", NULL};
    close(fd[0]);
    close(1);
    dup(fd[1]);
    close(fd[1]);
    exec("a.out", param);
} else {
    char *param[2] = {"b.out", NULL};
    close(fd[1]);
    close(0);
    dup(fd[0]);
    close(fd[0]);
    exec("b.out", param);
}
```

This entire code is part of the child process. We can implement pipe like behavior using the trick highlighted in red.